




IMPACT RANK 1 / ENTERPRISE AI SERVICES

The Frontier Labs Discovered the Hard Part: Implementation

OpenAI and Anthropic are moving beyond model access into deployment capacity – a signal that enterprise AI value will be captured where workflows, governance, and integration meet.



BUSINESS IMPACT

5/5

The strongest May signal: frontier AI is becoming implementation-led, especially for mid-market companies without deep internal AI teams.

For the first phase of generative AI, the market behaved as if distribution was the main problem: give employees access to a powerful model and productivity would follow. May's enterprise news points to a more mature reality. The bottleneck is not access. It is implementation.

Anthropic announced a new enterprise AI services company with Blackstone, Hellman & Friedman, Goldman Sachs, and other investors. Its stated target is mid-sized organizations that want Claude in core operations but do not have the in-house engineering depth to design, deploy, and maintain frontier workflows. TechCrunch also reported that OpenAI is preparing a similar services venture, reportedly called The Development Company, backed by major alternative asset managers.

The strategic logic is straightforward. Models are becoming more capable, but production value still depends on process mapping, data access, permissions, evaluation, change management, and human approval design. A bank, manufacturer, regional health group, or services firm rarely needs a blank chatbot. It needs an

applied engineering team that can sit with operators, identify where time disappears, and build around the tools people already use.

This is why the move resembles Palantir-style forward-deployed engineering more than conventional SaaS. The frontier labs and their investors are trying to create direct implementation channels into portfolio companies and mid-market operations. That makes enterprise AI less of a software subscription story and more of an operating transformation story.

For G-Nosis, the implication is encouraging. The market is validating the implementation-first thesis: diagnose workflows before selling automation, define governance before scaling agents, and prioritize measurable operational bottlenecks over generic AI enthusiasm.

What to watch: Newsletter angle: The next AI land grab is not chatbots. It is the ability to turn frontier models into governed operating capacity inside real companies.

Source note: Sources: Anthropic enterprise AI services announcement; TechCrunch enterprise JV reporting.

G-NOSIS READ

Build offers around workflow diagnosis, implementation, and governance – not generic “AI tool setup.”

NEWSLETTER USE


Use this as a one-page briefing article, then convert the central thesis into a short practical operator takeaway.



IMPACT RANK 2 / AGENTIC ADOPTION

AI Advantage Is Now Measured by Depth, Not Seat Count

OpenAI's B2B Signals research suggests the winning companies are not merely giving employees AI access – they are delegating deeper work to agents.



MORE INTELLIGENCE PER WORKER

3.5×

OpenAI says frontier firms now use 3.5× as much intelligence per worker as typical firms, up from 2× a year ago.

The first enterprise AI dashboard most leaders built was simple: seats deployed, logins, and message volume. Those numbers still matter, but they are no longer sufficient. OpenAI's B2B Signals report argues that the frontier of AI adoption is moving from access to depth.

According to the report, firms at the frontier now use 3.5 times as much intelligence per worker as typical firms. Message volume explains only 36% of that gap. The rest comes from richer context, more complex tasks, and more substantive outputs. In other words, leading companies are not just chatting more. They are asking AI to do more of the work.

The clearest marker is agentic usage. OpenAI reports that frontier firms send 16 times as many Codex messages per worker as typical firms. That matters because coding agents are an early proxy for delegated work: they inspect context, make changes, run tasks, and return work product. The same pattern is spreading into research, operations, finance, customer service, and internal tools.

This changes the executive question. Instead of asking, "How many people have AI?" leaders should ask, "Which workflows now run differently because AI is part of the operating loop?" A shallow deployment answers questions faster. A deep deployment redesigns handoffs, compresses cycle time, and lets humans approve higher-quality work instead of assembling it from scratch.

The adoption gap is likely to compound. Teams that learn how to delegate, evaluate, and govern agentic work will build operating muscle that competitors cannot copy simply by buying the same model subscription.

What to watch: Newsletter angle: Access is table stakes. The new AI maturity metric is how much meaningful work a company can safely delegate.

Source note: Source: OpenAI B2B Signals, May 2026.

G-NOSIS READ

Measure depth: workflows changed, hours saved, decisions accelerated, and agent tasks completed with human approval.

NEWSLETTER USE


Use this as a one-page briefing article, then convert the central thesis into a short practical operator takeaway.



IMPACT RANK 3 / FINANCE AGENTS

The CFO Stack Is Becoming Agentic

Finance is emerging as one of the first serious markets for production AI agents because its workflows are structured, expensive, measurable, and control-heavy.



CONTRACT THROUGHPUT SIGNAL

5x

OpenAI says its finance team used Codex to process 5x more contracts with the same-sized team.

Finance is where many AI promises become testable. The function has recurring workflows, clear controls, expensive professional labor, structured documents, and obvious deadlines. That makes it one of the strongest early markets for agentic automation.

OpenAI and PwC announced a collaboration focused on reimagining the office of the CFO with AI agents. The target workflows are not vague productivity tasks. They include planning, forecasting, reporting, procurement, payments, treasury, tax, accounting close, contract review, invoice review, and risk surfacing. These are operational rhythms that already define how finance teams run the business.

The important shift is from assistant to workflow participant. A finance agent can monitor payments and exceptions, compare invoices against policy, update a forecast when conditions change, prepare reporting materials, or surface risks before close. Human oversight remains essential, but the human starts from organized work product rather than a blank spreadsheet or inbox.

OpenAI is also positioning its own finance organization as customer zero. The company says its finance team used Codex to process five times more contracts with the same-sized team and used IR-GPT to manage more than 200 investor interactions during fundraising. Even allowing for vendor framing, the direction is clear: AI agents are being tested inside real finance operations, not just demo environments.

For G-Nosis, the connection to SAP and enterprise workflow automation is direct. Procurement, close, controls, vendor master data, exception handling, and reporting are exactly the kind of repeatable processes where AI agents can create practical value.

What to watch: Newsletter angle: Agents are not replacing CFOs. They are attacking the close, forecast, procurement queue, and exception backlog.

Source note: Source: OpenAI and PwC CFO collaboration announcement.

G-NOSIS READ

Prioritize finance workflows with clear rules, recurring volume, system context, and approval checkpoints.

NEWSLETTER USE

Use this as a one-page briefing article, then convert the central thesis into a short practical operator takeaway.



IMPACT RANK 4 / WORKFLOW TEMPLATES

The Winning AI Product May Be a Workflow Template

Anthropic's finance agent templates and Microsoft Office integrations show how AI moves from a blank chat box into the surfaces where professionals already work.



FINANCE TEMPLATES

10

Anthropic released ten ready-to-run templates covering pitchbooks, KYC, month-end close, valuation review, market research, and more.

A blank AI chat window is flexible, but flexibility can be a burden. Most professionals do not want to design an agent from scratch. They want a reliable workflow that understands the task, uses approved data, follows policy, and produces a familiar output.

Anthropic's May finance release points in that direction. The company announced ten ready-to-run agent templates for financial services: pitch builder, meeting preparer, earnings reviewer, model builder, market researcher, valuation reviewer, general ledger reconciler, month-end closer, statement auditor, and KYC screener. Each template packages instructions, domain knowledge, governed connectors, and subagents into a reference architecture that firms can adapt.

Just as important, Claude is being embedded across Microsoft Excel, PowerPoint, Word, and Outlook. That matters because professional work already lives in spreadsheets, decks, documents, and email. If an agent can build a model in Excel, carry the context into

PowerPoint, draft the client note in Outlook, and preserve auditability, adoption friction drops dramatically.

This is a more practical product shape than generic AI capability. Buyers understand named jobs. Analysts understand pitchbooks, close checklists, diligence Q&A, and KYC escalations. Compliance teams understand audit logs, permissions, and approval flows. Templates give all three groups a shared object to evaluate.

The lesson for builders is sharp: productize the workflow, not the model. The model is the engine; the template is the thing a business can buy, govern, and improve.

What to watch: Newsletter angle: The spreadsheet, deck, and inbox are becoming agent interfaces – and templates are the packaging layer.

Source note: Source: Anthropic finance agents and Microsoft 365 integration announcement.

G-NOSIS READ

Design offers as named workflow packages with inputs, outputs, approvals, and measurable operating metrics.

NEWSLETTER USE


Use this as a one-page briefing article, then convert the central thesis into a short practical operator takeaway.



IMPACT RANK 5 / AI INFRASTRUCTURE

AI Product Quality Depends on Invisible Infrastructure

Compute capacity, networking resilience, power, and regional deployment are no longer back-office concerns. They shape model availability, agent reliability, and enterprise trust.



GPU CAPACITY SIGNAL

220K+

Anthropic says its SpaceX partnership gives access to 220,000+ NVIDIA GPUs and 300+ MW of new capacity within the month.

AI users experience infrastructure as a product feature. A model feels useful when it is available, fast, accurate, and able to sustain long tasks. When usage limits tighten, latency rises, or agents fail mid-run, the problem may look like software. Underneath, it is often compute, networking, power, or regional capacity.

May made that infrastructure layer visible. Anthropic announced a SpaceX compute partnership that it says provides access to all capacity at the Colossus 1 data center – more than 300 megawatts and over 220,000 NVIDIA GPUs within the month. The same announcement included higher Claude Code and Claude Opus API limits. That is the market in miniature: more compute becomes more usable product.

OpenAI, meanwhile, released MRC – Multipath Reliable Connection – through the Open Compute Project, developed with AMD, Broadcom, Intel, Microsoft, and NVIDIA. The point is not marketing gloss. Large-scale model training depends on networks that move data between GPUs reliably. At frontier scale, congestion or a failed

link can stall expensive training jobs. Better networking helps models train faster and more predictably.

Google's recent AI recap emphasized eighth-generation TPUs designed for the agentic era. That phrase matters. Agents create long-running, tool-using, context-heavy workloads. They need more than raw model intelligence; they need dependable execution environments.

For enterprises, infrastructure also intersects with compliance. Regulated sectors care where inference runs, how data residency is handled, and whether capacity exists in-region. The AI winners will combine model quality with operational reliability.

What to watch: Newsletter angle: The AI race is a power, networking, and data-center race wearing a software costume.

Source note: Sources: Anthropic SpaceX capacity announcement; OpenAI MRC networking release; Google TPU/agentic infrastructure recap.

G-NOSIS READ

When evaluating AI platforms, assess rate limits, latency, regional capacity, auditability, and long-task reliability – not just benchmark scores.

NEWSLETTER USE

Use this as a one-page briefing article, then convert the central thesis into a short practical operator takeaway.